

Contents

Things to take note	4
FAQs	4
What is the sample average? Why is it a random variable?	4
What is the mean, variance and standard error of a Bernoulli random variable?	4
What is the sampling distribution?	4
What is the mean and variance of the sampling distribution? Derive them.	4
What is the Law of Large Numbers (LLN)?	5
What is the Central Limit Theorem (CLT)? What are its assumptions?	5
What does it mean when we say that \bar{Y} is an <i>estimator</i> of μ_Y ?	5
What does it mean for an estimator to be unbiased?	5
What does it mean for an estimator to be consistent?	5
What does it mean for an estimator to be efficient?	5
What does it mean when we say that \bar{Y} is the BLUE of μ_Y ?	5
What does it mean for an estimator to be a least squares estimator?	5
Prove that \bar{Y} is the least squares estimator of μ_Y	6
What is the t-statistic?	6
What is the p-value?	7
What is the confidence interval?	7
What is the sample covariance? What is its equation?	7
What is the difference between sample variance and the variance of the sample mean?	7
How do we differences-in-means? What is the t-statistic in a differences-in-means test?	9
Define the null and alternative hypothesis.	9
Under the null, what is the distribution of the test statistic?	9
Specify the significance level of the test, find critical values, and formulate the decision rule.	9
Calculate the actual value of the test statistic, t^{act}	9
Follow our decision rule and come to a conclusion.	9
What is an F-test? When do we do an F-test?	9
How do we do an F-test?	10
Show that the residual in the CEF decomposition is mean independent of X_i	10
What is OVB with a regression with more than one variable (or with one variable and a set of controls)	11
Under what circumstances does the LATE equal the TOT? Why? . . .	11
Time-series questions	11
What is the population autocorrelation?	11
What is the sample autocorrelation?	12
What does it mean for a time series to exhibit <i>strict stationarity</i> ? . . .	12

What does it mean for a time series to exhibit <i>weak stationarity</i> ? . . .	12
What is an AR(n) model?	12
What is the difference between an AR(n) model and an AR(n) process? . . .	12
How do we solve the AR(1) process?	12
What is the first moment, second moment and autocorrelation of the AR(1) process?	13
What are the sampling properties of OLS?	13
What is the difference between a predicted value and a forecast? . . .	13
What's the difference between $\hat{Y}_{t t-1}$ and $Y_{t t-1}$?	13
What is RMSFE? Define it, give the equation.	14
What is Granger causality?	14
Worked example: Does unemployment Granger-cause in- flation?	15
What does it mean for a time series to exhibit a <i>deterministic trend</i> ? . .	15
What does it mean for a time series to be trend stationary?	15
What does it mean for a time series to exhibit a <i>stochastic trend</i> ? . . .	15
What is the equation of a random walk with drift?	16
What are the mean, variance and covariance of a random walk with drift? .	16
How do we detrend a deterministic trend?	16
How do we detrend a stochastic trend?	16
What is order of integration?	16
What is a unit root? What issues arise when we have a unit root? . .	16
How do we test for a stochastic trend/unit root? Be explicit about the procedure	17
What's the difference between the Dickey-Fuller, the DF with trend, and the ADF? When do we use what?	17
What's the problem with a break?	17
How do we test for a break when the break date is known?	18
How do we test for a break when the break date is unknown?	18
How do we test for a break when there can be multiple breaks?	18
What's the Chow test?	18
What's the QLR test?	18
What is cointegration?	18
What is the cointegrating coefficient?	18
How do we test for cointegration when the cointegrating coefficient is known?	18
How do we test for cointegration when the cointegrating coefficient is unknown?	18
What is h-step ahead RMSFE? Derive it for $h = 1, 2, 3, 4$	19
Steps for hypothesis testing	19
How to actually run a hypothesis test	19
Notes on hypothesis testing	19

Hypothesis testing on regression parameters	20
p-values	20
Testing the hypothesis that one sample mean is greater than another.	20
Regression analysis and interpretation	20
Flowchart for regression testing	20
I added a new variable in my regression. Should I expect the standard error on my coefficient of interest to go up or down?	20
Should I include a new variable in my regression?	21
What is the formula for omitted variable bias (OVB) for a regression with more than one variable?	22
How do we test an instrument for exogeneity?	23
How do we test an instrument for exclusion?	23
What's the difference between exclusion and exogeneity?	23
How do we test an instrument for relevance?	24
Heterogeneity	25
Selection bias	25
R stuff	25
F-test	25
Instrument exogeneity	26
Instrumental variables estimation in R	26
ILS	26
Manual 2SLS	26
Automated 2SLS	26

Things to take note

Always give economic intuition explanation especially for regression interpretation questions. Put on your PolSoc hat.

When they ask for interpretation of the coefficient or whatever: don't just talk about the straightforward interpretation, see if you can talk more about it. Is it (plausibly) the LATE? TOT? ATE of compliers?

When they ask for internal validity: check random assignment (exogeneity) and relevance.

When they ask for external validity: check how close the group under study is to

FAQs

What is the sample average? Why is it a random variable?

Sample average is \bar{Y} . It is a random variable because it is a function of random variables of the population.

What is the mean, variance and standard error of a Bernoulli random variable?

Let \hat{p} be the sample mean (equivalently written as \bar{X}).

$$\begin{aligned}E\hat{p} &= p \\var(X_i) &= p(1-p) \\var(\hat{p}) &= p(1-p)/n \\se(\hat{p}) &= \frac{\hat{sd}(\hat{p})}{\sqrt{n}} = \left(\frac{\hat{p}(1-\hat{p})}{n}\right)^{1/2}\end{aligned}$$

What is the sampling distribution?

The distribution of \bar{Y} .

What is the mean and variance of the sampling distribution? Derive them.

$$\begin{aligned}E(\bar{Y}) &= E(1/n \sum_{i=1}^N Y_i) = 1/n \sum_{i=1}^N EY_i = E[Y](i.i.d) = \mu_Y \\var(\bar{Y}) &= var\left(\frac{1}{n} \sum_{i=1}^N Y_i\right) = \frac{1}{n^2} var\left(\sum_{i=1}^N Y_i\right) = \sigma^2/n \quad (Y_i \perp Y_j)\end{aligned}$$

What is the Law of Large Numbers (LLN)?

If Y_i are i.i.d with $E(Y_i) = \mu_Y$ and $var(Y_i) = \sigma_Y^2 < \infty$ then

$$\bar{Y} \xrightarrow{p} \mu_Y.$$

What is the Central Limit Theorem (CLT)? What are its assumptions?

Assumptions: Y s must be i.i.d, $0 < var(Y_i) < \infty$.

As $n \rightarrow \infty$, the distribution of

$$\frac{\bar{Y} - \mu_Y}{\sigma_{\bar{Y}}} \sim N(0, 1)$$

What does it mean when we say that \bar{Y} is an *estimator* of μ_Y ?

6. An estimator is a random variable that is a function of a sample of data drawn randomly from the population.

What does it mean for an estimator to be unbiased?

7. An estimator \hat{a} is a consistent estimator of a iff $E(\hat{a}) = a$.

What does it mean for an estimator to be consistent?

8. \hat{a} is a consistent estimator of a if as N gets large, for any $\epsilon > 0$, the probability that $\hat{a} - a < \epsilon$ tends to zero.

What does it mean for an estimator to be efficient?

9. An efficient estimator is an estimator that has low variance.

What does it mean when we say that \bar{Y} is the BLUE of μ_Y ?

The Best Linear Unbiased Estimator (BLUE) is the estimator that has the smallest variance.

What does it mean for an estimator to be a least squares estimator?

An estimator m minimises the sum of squared differences between the observations of the sample and m .

Prove that \bar{Y} is the least squares estimator of μ_Y .

12. Lecture 4 slide 18/20.

What is the t-statistic?

The t-statistic is any statistic of the form

$$t = \frac{\hat{\beta} - \beta}{se(\hat{\beta})}$$

where $se(\hat{\beta})$ is the standard error of the estimated parameter. Note the difference between that and $\hat{se}(\beta)$.

The former is the standard error of the estimated parameter, which is something like the square root of the *variance of the sample mean*.

The latter is the square root of the *sample variance*.

We have the following (replace “sample mean” with “estimated parameter”):

- $Var(X) = \sigma^2$ is the population variance.
- $sd(X) = \sigma$ is the population standard deviation.
- $\hat{Var}(X) = s^2$ is the sample variance.
- $\hat{sd}(X) = \sqrt{\hat{Var}(X)} = s$ is the sample standard deviation.
- $Var(\hat{X})$ is the variance of the sample mean.
- $sd(\hat{X})$ is the standard deviation of the sample mean.
- $se(\hat{X})$ is the standard error of the sample mean. It estimates the standard deviation of the sample mean, which is unknown.

The relationship between them is the following:

The sample variance is an unbiased estimator of the population variance. That is,

$$\hat{Var}(X) \equiv s^2 \rightarrow \sigma^2.$$

The variance of the sample mean is equal to the population variance divided by n .

$$Var(\hat{X}) = \frac{\sigma^2}{n}.$$

This allows us to write the following:

$$se(\hat{X}) \rightarrow sd(\hat{X}) = \sqrt{Var(\hat{X})} = \sqrt{\frac{\hat{Var}(X)}{n}} = \hat{sd}(X)/\sqrt{n}$$

What is the p-value?

p-value or probability value is the probability of obtaining test results at least as extreme as the results actually observed (t^{act}) during the test, assuming that the null hypothesis is correct.

Entirely equivalently, the p-value is the lowest significance level under which the null hypothesis would be rejected.

What is the confidence interval?

An X% two-sided confidence interval for μ_Y is a random interval that contains the true value of μ_Y X% of the time. Given the sample average we observe in our randomly drawn sample, there is a 95% chance that the true population mean lies in the interval between A and B. Note that this is a property of the CLT — sample means must follow a normal distribution, so we can make claims about what the population mean *should* be.

- 90% confidence interval is $\pm 1.64SE$
- 95% confidence interval is $\pm 1.96SE$
- 99% confidence interval is $\pm 2.58SE$

What is the sample covariance? What is its equation?

The sample covariance is the sample analogue of the population covariance. It is

$$\hat{Cov}(X, Y) = \sum_{i=0}^N (X_i - \bar{X})(Y_i - \bar{Y})$$

What is the difference between sample variance and the variance of the sample mean?

The sample mean \bar{Y} is a random variable (it is after all a function of random variables Y_i) and, being a random variable, it has a mean $E[\bar{Y}]$ and a variance $\sigma_{\bar{Y}}^2$. It can be shown that $E[\bar{Y}] = \mu_Y$ and $Var(\bar{Y}) \equiv \sigma_{\bar{Y}}^2 = \sigma_Y^2/n$.

But we don't know σ_Y , how can we know $\sigma_{\bar{Y}}$? We need to estimate the variance of the sample mean. It turns out that we estimate the *variance of the sample mean* $\sigma_{\bar{Y}}^2$ with what is called the *sample variance*, s_Y^2 .

$$\sigma_{\bar{Y}} \equiv SD(\bar{Y})$$

This can be estimated by the *sample variance*. s_Y^2 , the sample variance, is a random variable, and is a consistent estimator of the variance of the sample and the population variance.

$$\hat{sd}(X) = s_Y^2 = \frac{1}{n-1} \sum_i^n (Y_i - \bar{Y})^2 \rightarrow \sigma_Y^2$$

There is also the term “Standard error of \bar{Y} ”, or $SE(\bar{Y})$: this is equivalent to $\hat{\sigma}_{\bar{Y}}$. The notation is a bit confusing but I believe

$$\begin{aligned}\hat{\sigma}_{\bar{Y}} &\equiv SE(\bar{Y}) \\ SE(\bar{Y}) &\equiv \hat{\sigma}_{\bar{Y}} \rightarrow \sigma_{\bar{Y}},\end{aligned}$$

that is to say that the standard error of \bar{Y} is an estimator of the standard deviation of \bar{Y} . From the previous two equations we can write

$$SE(\bar{Y}) \equiv \hat{\sigma}_{\bar{Y}} = \frac{s_Y}{\sqrt{n}}$$

Not sure what the relationship between all of these things is. I think standard error of \bar{Y} , $SE(\bar{Y})$, is another way to say the sample standard deviation s_Y divided by \sqrt{n} , which is an (unbiased and consistent?) estimator of the *standard deviation of the sample mean* $\sigma_{\bar{Y}}$. But why must there be two different terms for the same fucking thing?

So when we are normalising

$$\frac{(\bar{Y} - \mu_Y)}{\sigma_{\bar{Y}}}$$

we can simply write

$$\frac{(\bar{Y} - \mu_Y)}{SE(\bar{Y})}$$

Similarly, when doing difference-in-means tests, we can write

$$\frac{(\bar{Y}_a - \bar{Y}_b)}{SE(\bar{Y}_a - \bar{Y}_b)}$$

which equals

$$SE(\bar{Y}_a - \bar{Y}_b) = \sqrt{\frac{s_a^2}{n_a} + \frac{s_b^2}{n_b}}$$

by the variance rules $Var(A - B) = Var(A) - 2Cov(A, B) + Var(B)$ and here $Cov(A, B) = 0$ by the fact that A and B are independent samples from different populations.

How do we differences-in-means? What is the t-statistic in a differences-in-means test?

Define the null and alternative hypothesis.

$$H_0 : \mu_w = \mu_m$$

$$H_1 : \mu_w \neq \mu_m$$

Under the null, what is the distribution of the test statistic?

The t-statistic for testing differences in means is

$$t = \frac{\bar{Y}_m - \bar{Y}_w}{SE(\bar{Y}_m - \bar{Y}_w)}$$

When n_m and n_w are large, then by the CLT, the t-statistic has a standard normal distribution when the null hypothesis is true.

Specify the significance level of the test, find critical values, and formulate the decision rule.

Suppose we wanted to test the hypothesis at a 5% significance level.

Under the null hypothesis, the distribution of the test statistic is approximately standard normal.

At the 5% significance level, the critical value $c_\alpha = 1.96$.

Decision rule: reject H_0 if $|c_\alpha| > 1.96$.

Calculate the actual value of the test statistic, t^{act} .

The standard error can be calculated as

$$SE(\bar{Y}_m - \bar{Y}_w) = \sqrt{s_m^2/n_m + s_w^2/n_w}$$

Substitute this value into the t-statistic and find t^{act} .

Follow our decision rule and come to a conclusion.

Given that $t^{act} > 1.96$, we reject the null that the means are equal at the 5% level.

What is an F-test? When do we do an F-test?

- Checking if regression coefficients are significantly different from zero

How do we do an F-test?

Show that the residual in the CEF decomposition is mean independent of X_i .

1.
(a) [50%] Show that the residual e_i in the identity $Y_i = E[Y_i|X_i] + e_i$ is mean independent of X_i .

Claim: $E(e_i | X_i) = E(e_i)$.

We have the identity $Y_i = E[Y_i|X_i] + e_i$ with $E(e_i) = 0$

Taking conditional expectations on both sides and by construction,

by linearity, we have:

$$E[Y_i|X_i] = E[E[Y_i|X_i] + e_i | X_i] \\ \text{By LIE, } E[E(\cdot)|X] = E(\cdot), \text{ so}$$

$$= E[Y_i|X_i] + E[e_i|X_i]$$

$$\Rightarrow E[e_i|X_i] = 0$$

We know by construction that $E(e_i) = 0$ in the CEF decomposition identity, so

$$E(e_i|X_i) = E(e_i) = 0. \text{ Shown } \square.$$

What is OVB with a regression with more than one variable (or with one variable and a set of controls)

Great Job.

OVB for a regression with more than one variable (say X_2 is

$$\text{Short: } \alpha_0 + \beta_0 X_1 + \gamma_0 Z + e$$

$$\text{Long: } \alpha_1 + \beta_1 X_1 + \beta_2 X_2 + \gamma_1 Z + u$$

omitted but Z is
a control/set of
controls

Then

$$\begin{aligned} \beta_1 &= \frac{\text{Cov}(Y, \tilde{X}_1)}{\text{Var}(\tilde{X}_1)} \quad \text{where } X_1 = \delta_0 + \delta_1 Z + \tilde{X}_1 \\ &= \frac{\text{Cov}(\tilde{\alpha}_1 + \beta_1 X_1 + \beta_2 X_2 + \gamma_1 Z, \tilde{X}_1)}{\text{Var}(\tilde{X}_1)} \quad \left. \begin{array}{l} \text{by constraints} \\ \text{if } \tilde{X}_1 = 0 \\ \text{if } \tilde{X}_1 Z = 0 \end{array} \right\} \\ &= \frac{\text{Cov}(\tilde{\alpha}_1 + \beta_1 X_1 + \beta_2 X_2 + \gamma_1 Z, \tilde{X}_1)}{\text{Var}(\tilde{X}_1)} \\ &= \frac{\text{Cov}(\beta_1 X_1, \tilde{X}_1)}{\text{Var}(\tilde{X}_1)} + \frac{\text{Cov}(\beta_2 X_2, \tilde{X}_1)}{\text{Var}(\tilde{X}_1)} \\ &= \beta_1 + \beta_2 \frac{\text{Cov}(X_2, \tilde{X}_1)}{\text{Var}(\tilde{X}_1)} \quad \square \end{aligned}$$

Say X_2 is omitted but there is a set of controls Z .

**Under what circumstances does the LATE equal the TOT?
Why?**

Only compliers and defiers affect the LATE, because always- and never-takers always work in a deterministic way. Without defiers, LATE reduces to the average treatment effect of compliers, and without always-takers, those who are taking the treatment are the ones who have been offered treatment, so the LATE recovers the TOT.

Time-series questions

What is the population autocorrelation?

The n th population autocovariance is the correlation between Y and its n th lag.

$$\text{corr}(Y_t, Y_{t-j})$$

What is the sample autocorrelation?

The sample autocorrelation is the sample version of the pop autocorrelation:

$$\hat{corr}(Y_t, Y_{t-j}) = \frac{1}{T-j-1} \sum_{t=j+1}^T (Y_t - \bar{Y}_{j+1,T})(Y_{t-j} - \bar{Y}_{T-j})$$

Why are the subscripts what they are? Let's use some numerical examples to clear things up. Consider taking the 2nd sample autocorrelation: that is, the sample correlation of Y_t with Y_{t-2} .

The first sample $\bar{Y}_{j+1,T}$ is the sample mean, counting from $t = 3$ to $t = T$. And the second sample $\bar{Y}_{1,T-j}$ is the sample mean, counting from $t = 1$ to $t = T - j$. This makes sense because when you are doing the autocorrelation, you can't start measuring autocorrelations until you have enough lags (in this case, $t=3$).

What does it mean for a time series to exhibit *strict stationarity*?

A time series Y_t exhibits strict stationarity iff its distribution does not change over time.

What does it mean for a time series to exhibit *weak stationarity*?

A time series Y_t exhibits weak stationarity iff its first and second moments (mean, variance and autocovariance) *exist* and are *constant over time*. They must be finite.

What is an AR(n) model?

Autoregressive model: a model with (up to) n lags of itself

$$Y_t = \beta_0 + \beta_1 Y_{t-1}$$

What is the difference between an AR(n) model and an AR(n) process?

How do we solve the AR(1) process?

We can solve the AR(1) process by backward substitution. Note that Y_{t-1} can itself be written as an AR(1) process, so keep substituting in until we get LHS y_t and RHS y_0 .

What is the first moment, second moment and autocorrelation of the AR(1) process?

If $\beta_1 < 1$ (no unit root), and $\text{var}(u_t) = \sigma^2$ (constant),

$$E(Y_0) = \frac{\beta_0}{1 - \beta_1}$$
$$\text{var}(Y_0) = \frac{\sigma^2}{1 - \beta_1^2}$$

and we can derive the ACF as

$$\rho_j = \text{corr}(Y_t, Y_{t-j}) = \beta_1^j$$

in a similar backwards substitution process.

What are the sampling properties of OLS?

What allows us to estimate consistently the coefficient on β_l ? In the regular OLS regression we require i.i.d. But a similar result holds for non-i.i.d data provided they are weakly stationary and “weakly dependent”: that is, the nth autocovariance $\rho_n \rightarrow 0$ as n tends to infinity.

This is precisely fulfilled when we don’t have a unit root in the AR(1) model, because $\rho_j = \text{corr}(Y_t, Y_{t-j}) = \beta_1^j$, and anything < 1 taken to a power tends to 0.

In the AR(1) model, the idea is that the influence of $Y(t-j)$ on Y_t is going to be very small because $\beta_1 < 1$ and thus that raised to a power j is going to tend to 0 as t gets large.

What is the difference between a predicted value and a forecast?

predicted values are in-sample, forecast are out-of-sample

What’s the difference between $\hat{Y}_{t|t-1}$ and $Y_{t|t-1}$?

$Y_{t|t-1}$ is the forecast given all the data from $t = 0$ to $t - 1$ using the population (true unknown) coefficients.

$\hat{Y}_{t|t-1}$ is the “sample” forecast using the coefficients $\hat{\beta}_0$, $\hat{\beta}_1$ that were estimated in the OLS regression.

What is RMSFE? Define it, give the equation.

The one-period ahead forecast error is

$$Y_t - \hat{Y}_{t|t-1},$$

that is to say, the difference between the actual out-of-sample value Y_t and the predicted value

Root-mean-squared-forecast-error

The RMSFE is

$$\sqrt{E[Y_t - \hat{Y}_{t|t-1}]},$$

a measure of the magnitude of the typical forecasting “mistake”.

If we look at the error here,

$$Y_t - \hat{Y}_{t|t-1},$$

it can be decomposed into the genuinely unforecastable error (random shocks), and the forecast error due to estimation error of our coefficients. That is to say,

$$Y_t - \hat{Y}_{t|t-1} = (Y_t - \hat{Y}_{t|t-1}) + (\hat{Y}_{t|t-1} - \hat{Y}_{t|t-1}) = u_t + (\beta_0 + \hat{\beta}_0) + (\beta_1 + \hat{\beta}_1)Y_{t-1}$$

The bigger our sample, the lower the estimation error will become, but the genuinely unforecastable error will not decrease.

What is Granger causality?

X_t Granger-causes Y_t if including lags of X_t helps to predict Y_t over and above just lags of Y_t .

$$E(Y_t|Y_{t-1}, Y_{t-2} \dots X_{t-1}, X_{t-2}) \neq E(Y_t|Y_{t-1}, Y_{t-2} \dots)$$

The Granger causality statistic is the F-statistic testing the hypothesis that the coefficient on all the values of one of the variables are zero. This implies that the regressors have no predictive content for Y_t beyond that contained in the other regressors.

Worked example: Does unemployment Granger-cause inflation?

If we have an ADL(1,1) model

$$\Delta Inf_t = \beta_0 + \beta_1 \Delta Inf_{t-1} + \delta_1 Unrate_{t-1} + \delta_2 Unrate_{t+2} + u_t$$

we test whether lags of Unrate are significant with the following F-test:

$$H_0 : \delta_1 = \delta_2 = 0$$
$$H_1 : \delta_1 \neq 0 \quad \text{or} \quad \delta_2 \neq 0$$

What does it mean for a time series to exhibit a *deterministic trend*?

A time series exhibits a deterministic trend if it has a trend that is a deterministic function of time:

$$Y_t = \alpha t + \beta_0 + \beta_1 Y_{t-1}$$

where α is some constant.

What does it mean for a time series to be trend stationary?

If it exhibits stationary deviations from a deterministic trend (i.e. once you remove the deterministic trend it becomes stationary)

What does it mean for a time series to exhibit a *stochastic trend*?

Basically just a random walk (or a random walk with trend)

$$Y_t = Y_{t-1} + u_t$$

or

$$Y_t = \alpha_1 t + Y_{t-1} + u_t$$

Solving backwards we obtain

$$Y_t = \alpha_1 t + \sum_{j=1}^t u_j$$

What is the equation of a random walk with drift?

$$Y_t = \alpha_1 + Y_{t-1} + u_t$$

What are the mean, variance and covariance of a random walk with drift?

We have the random walk with drift as

$$Y_t = \alpha_1 t + \sum_{j=1}^t u_j$$

Assuming that u_j is *i.i.d* with distribution $(0, \sigma^2)$,

$$E(Y_t) = \alpha_1 t,$$

$$var(Y_t) = var\left(\sum_{j=1}^t u_j\right) = t\sigma^2,$$

$$cov(Y_t, Y_{t-j}) = (t-j)\sigma^2$$

How do we detrend a deterministic trend?

Regress Y_t on a deterministic function of time and take the residuals.

$$Y_t = \alpha_0 + \alpha_1 t$$

How do we detrend a stochastic trend?

Take first differences.

What is order of integration?

We say that Y_t is integrated of order x if Y_t must be differenced x times to remove its stochastic trend.

What is a unit root? What issues arise when we have a unit root?

A unit root is a stochastic trend.

$$Y_t = \beta_0 + \beta_1 Y_{t-1} + u_t$$

where $\beta_1 = 1$.

There are two issues with a unit root:

$$E(\hat{B}_1^{OLS}) = 1 - \frac{5.3}{T}$$

Firstly, the distribution of the OLS estimator and the t-statistic is not normal even in large samples. We can't use normal critical values, we will get a biased OLS estimate.

Secondly, you get *spurious regression*: stochastic trends can make two unrelated time series appear related. Stochastically trending processes will tend to correlate with any other process that exhibits a trend. We will spuriously reject the null of no relationship as sample increases.

How do we test for a stochastic trend/unit root? Be explicit about the procedure

Do an F-test: subtract Y_{t-1} from both sides of an AR(1) model.

$$\Delta Y_t = \beta_0 + \delta Y_{t-1} + u_t,$$

where $\delta = \beta_1 - 1$

Then test the null hypothesis with an F-test with $H_0 = \delta = 0$

Use the Dickey-Fuller critical values, *not* the normal critical values.

What's the difference between the Dickey-Fuller, the DF with trend, and the ADF? When do we use what?

- Dickey-Fuller: standard unit test
- DF with trend: unit test with deterministic trend
- ADF: Augment the DF regression model with lags of ΔY_t in the RHS:

$$\Delta Y_t = \beta_0 + \alpha t + \delta Y_{t-1} + \sum_{j=1}^p \delta_j \Delta Y_{t-j} + u_t$$

What's the problem with a break?

They cause in-sample estimates of coefficients to be biased and destroy the external validity of time series models

How do we test for a break when the break date is known?

Chow test: just a standard F test. Have a dummy variable 0 before the break, 1 after the break.

Estimate

$$Y_t = \beta_0 + \beta_1 X_t + \gamma_0 D_{\tau,t} + \gamma_1 D_{\tau,t} X_t + u_t$$

Test the null hypothesis that $H_0 : \gamma_0 = \gamma_1 = 0$.

How do we test for a break when the break date is unknown?

QLR test: do the Chow test for multiple breaks and take the maximum critical values

How do we test for a break when there can be multiple breaks?

You can't

What's the Chow test?

What's the QLR test?

What is cointegration?

When two variables are related by some common constant:

$$Y_t = \theta X_t$$

What is the cointegrating coefficient?

The θ in $Y_t = \theta X_t$

How do we test for cointegration when the cointegrating coefficient is known?

[TODO]

How do we test for cointegration when the cointegrating coefficient is unknown?

[TODO]

If we know that Y and X are cointegrated, then by taking first difference of Y_t , we have removed the spurious regression.

The idea here is that if δY_t is positive, we subtract a bit (λ) from Y_{t-1} to “correct” for this

The parameter λ tells us how much Y_t adjusts to disturbances in eqm

You have the short-run relationship (which is just the differences) and the long run relationship

Where α and β are known, estimate these with Engel and Granger.

What is h-step ahead RMSFE? Derive it for $h = 1, 2, 3, 4$.

Steps for hypothesis testing

1. Define the null and alternative hypothesis.
2. Under the null hypothesis, what is the distribution of the test statistic?
3. Specify significance levels, calculate confidence intervals and critical values.
4. Come up with a decision rule: “Reject H_0 if $t^{act} > c_\alpha$ ”
5. Calculate the actual value of the test statistic from the data.
6. Reject the null hypothesis if the t-statistic is larger than the critical value.

How to actually run a hypothesis test

Notes on hypothesis testing

In order to use the CLT,

$$\frac{\bar{Y} - \mu_Y}{\sigma_{\bar{Y}}} \sim N(0, 1)$$

you need to derive that $E(\bar{Y}) = \mu_Y$ and $Var(\bar{Y}) = \frac{\sigma_Y^2}{n}$ first.

The sample average \bar{Y} is randomly distributed with $N(\mu_Y, \sigma_Y^2/n)$.

The t-statistic is a random variable. It is given by

$$t = \frac{\bar{Y} - \mu_Y}{SE(\bar{Y})}$$

The actual calculated test statistic, t^{act} , is just a number that you get when you plug all of that in.

Hypothesis testing on regression parameters

$$\omega_{\beta_1} = \frac{E(X_i - EX_i)^2 u_i^2}{[E(X_i - EX_i)^2]^2}$$

$$sd(\hat{\beta}_1) = \frac{1}{\sqrt{n}} \omega_{\beta_1}$$

$$se(\hat{\beta}_1) = \frac{1}{\sqrt{n}} \hat{\omega}_{\beta_1}$$

p-values

Testing the hypothesis that one sample mean is greater than another.

Remember that the standard errors must be added together.

Regression analysis and interpretation

Flowchart for regression testing

Check flowchart PNG

I added a new variable in my regression. Should I expect the standard error on my coefficient of interest to go up or down?

TLDR: It depends on the covariances. On the one hand, the new variable will explain somewhat the dependent variable, which causes standard error to go down; on the other hand, ...

Assuming homoskedasticity, the standard error for $\hat{\beta}_1$ can be written as

$$se(\hat{\beta}_1) = \left(\frac{1}{n} \frac{\hat{var}(u)}{\hat{var}(\tilde{X}_1)} \right)^{\frac{1}{2}}$$

where \tilde{X}_1 is the residual from an OLS regression of X_1 on (X_2, \dots, X_k) .

Let's use an example to make things clearer. Suppose you had the regression

$$Wages = \beta_0 + \beta_1 Experience + u$$

Now we want to add a new variable, gender, to the regression to get the “long” regression:

$$Wages = \gamma_0 + \gamma_1 Experience + \gamma_2 Gender + v$$

What will happen to the standard error of the coefficient? Well, gender should explain wages to some extent, so we expect $var(u)$ to go down. On the other hand, $\widetilde{Experience}_1$ is the residual from an OLS regression of Experience on Gender:

$$Experience = \pi_0 + \pi_1 Gender + \widetilde{Experience}_1$$

and given that some of experience is explained by gender (maybe men have more years in the workforce in general — no break for childbearing), $var(\widetilde{X}_1)$ will decrease.

So the overall effect is ambiguous. But in general, if

$$Cov(Gender, Wages) > Cov(Gender, Experience),$$

$Cov(Gender, Experience)$ is π_1

$Cov(Gender, Wages)$ is γ_2

that is, gender explains wages more than it explains experience, then the standard error will go down when adding more regressors.

The intuition is that the former is the decrease in the residual u of the short regression: the higher $Cov(Gender, Wages)$ is, the more of wages is explained, and the lower u falls.

The latter is the decrease in the variance of experience after it has been explained by gender. The higher $Cov(Gender, Experience)$ is, the more gender is explained by experience, and the smaller the residual $\widetilde{X}_{\text{whatever}}$ is going to be.

Should I include a new variable in my regression?

TLDR: Not if your new factor is possibly endogenous, because that will cause all the other variables to be estimated with error. You should *only* add new variables if they are exogenous with the error term!

Suppose you had the regression

$$Wages = \beta_0 + \beta_2 Gender + u$$

and we had good reason to believe that gender was exogenous with the error term (OR); that is, $Cov(Gender, u) = 0$.

Now suppose someone suggests that you add occupation into the regression to get a “cleaner estimate” of the wage effect. That is,

$$Wages = \gamma_0 + \gamma_1 Gender + \gamma_2 Occupation + v$$

It's true that if occupation was exogenous this would indeed increase the precision of the estimate. This is because

$$se(\hat{\beta}_1) = \frac{\hat{var}(u)}{\hat{var}(\tilde{x})}$$

and given that occupation explains some part of wages, and is exogenous with gender (what we assumed), only $var(u)$ will go down, $var(\tilde{x})$ will remain unchanged.

But this is only if $Cov(Occupation, v) = 0$! If occupation is endogenous (for instance, if ability determines occupation and wages), then this would cause the OLS estimates of all the variables (including the coeff on gender which was previously the correct causal interpretation) to be wrongly estimated.

What is the formula for omitted variable bias (OVB) for a regression with more than one variable?

Set up the “short” and “long” regressions, and substitute the long regression into the short regression giving you β_1 :

$$\beta_1 = \beta_1 + \beta_2 \frac{Cov(X_2, \tilde{X}_1)}{Var(\tilde{X}_1)}$$

Great Job.

OVB for a regression with more than one variable (say X_2 is omitted but Z is a control / set of controls)

Short: $Y_i = \beta_0 + \beta_1 X_i + \beta_2 X_2 + \beta_3 Z + e_i$

Long: $Y_i = \alpha_0 + \beta_1 X_i + \beta_2 X_2 + \beta_3 Z + u_i$

Then

$$\beta_1 = \frac{\text{Cov}(Y, \tilde{X}_1)}{\text{Var}(\tilde{X}_1)}$$

where $X_i = \delta_0 + \delta_1 Z + \tilde{X}_i$

$$= \frac{\text{Cov}(\delta_0 + \delta_1 Z + \tilde{X}_i, \tilde{X}_i)}{\text{Var}(\tilde{X}_i)}$$

$\left\{ \begin{array}{l} E\tilde{X}_i = 0 \\ E\tilde{X}_i Z = 0 \end{array} \right\}$ by construction

$$= \frac{\text{Cov}(\delta_1 Z, \tilde{X}_i)}{\text{Var}(\tilde{X}_i)} + \frac{\text{Cov}(\tilde{X}_i, \tilde{X}_i)}{\text{Var}(\tilde{X}_i)}$$

$$= \delta_1 + \beta_2 \frac{\text{Cov}(X_2, \tilde{X}_1)}{\text{Var}(\tilde{X}_1)}$$

How do we test an instrument for exogeneity?

We run an F-test on the coefficients of the regressions against the residuals.

$$\hat{u} = Y - \hat{\beta}_0 + \hat{\beta}_1 X$$

and we check if $\beta_1 = 0$. We can never prove that an instrument is exogenous. We can only fail to reject the null of exogeneity.

How do we test an instrument for exclusion?

You can't! This is a story about the causal model you have to tell.

What's the difference between exclusion and exogeneity?

Definition of exclusion: consider a (possibly endogenous) variable X and a proposed instrumental variable Z . In a causal model of Y on X and Z , the coefficient on Z should be zero: that is, Z has no effect on Y other than through X .

Angrist (1990) gives an example of an exogenous instrumental variable that was not endogenous. Angrist wanted to find the effect of serving in the military on wages. So he would like to run the following regression:

$$Wage = \beta_0 + \beta_1 MilSvc + u$$

But of course military service is endogenous with the error term here. So instead he used the fact that people were drawn in a lottery to be drafted for the Vietnam War. Is this IV exogenous? Surely so. It was randomly assigned i.e. it can't be correlated with anything.

But does it satisfy *exclusion*? In fact, no. Because of the fact that you couldn't be drafted if you were in school, people who got picked to join might have stayed in school longer, continuing further study, which would have an effect on their wages. So in the following causal model

$$Wage = \delta_0 + \delta_1 MilSvc + \delta_2 Lottery + v$$

$\delta_2 \neq 0$ and thus exclusion is not satisfied.

How do we test an instrument for relevance?

Suppose we have the “short” regression

$$Y = \beta_0 + \beta_1 X + \beta_2 Z + u$$

where Z is a vector of control variables, and X is possibly endogenous. We wish to instrument X with D and so we run the following first-stage regression:

$$X = \gamma_0 + \gamma_1 D + \gamma_2 Z + v$$

We set up the following hypothesis test:

$$H_0 : \gamma_1 = 0$$

$$H_1 : \gamma_1 \neq 0$$

and we do an F-test by looking at the sum of squared residuals in the *restricted model* setting $\gamma = 0$ and in the *unrestricted model* (that ILS regression):

$$LR = \frac{SSR_{rs} - SSR_{un}}{SSR_{un}/(n - k - 1)} \rightarrow \chi_q^2$$

If the F-statistic is sufficiently greater than 0, than the instrument is relevant; if the F-statistic is greater than 10, than the instrument is relevant.

Heterogeneity

We are interested in the causal effect of X on Y , and the magnitude of that effect is β_1 .

The key additional assumption to make in the case of heterogeneity is that

$$E[\beta_{1i}|X_i] = E\beta_{1i},$$

that is to say, that the average causal effect of the treatment does not vary systematically with the treatment. For instance, if X_i was a skills learning program and smarter people were more likely to be offered the treatment $X_i = 1$, then

$$E[\beta_1|X_i = 1] > E[\beta_1|X_i = 0]$$

This mean independence assumption is usually stated as a stronger independence assumption: both β_{0i} and β_{1i} are independent of X_i .

Selection bias

Selection bias is the difference in the untreated outcomes between people who were treated, and people who were not:

Among the people who were treated, what would be their potential outcomes if they were not treated? That is to say, what is $Y_{0i}|D_i = 1$, or entirely equivalently, $E[\beta_{0i}|D_i = 1]$?

And among the people who were *not* treated, what is their outcome? That is to say, what is $Y_{0i}|D_i = 0$, or entirely equivalently, $E[\beta_{0i}|D_i = 0]$?

The difference between these two groups is selection bias. Going back to the running example, if you choose only smart people to participate in your skills learning program, then selection bias would be positive.

When the independence assumption fails, the OLS regression consistently estimates the TOT + SB where TOT is the treatment on the treated and the SB is the selection bias term.

R stuff

F-test

```
linearHypothesis(model, matchCoefs(model, "regex_string_of_coefs_to_match"),  
test = "F")
```

Instrument exogeneity

Slide 123

1. Compute the 2SLS residuals
2. Perform a homoskedasticity-only F-test for the null that the coeffs on the instruments are null

Instrumental variables estimation in R

Suppose we want to estimate this model

$\text{lwage} = \beta_0 + \beta_1 \text{educ} + \{\text{demog}, \text{family}\} + u$

using `nearc4` as an instrumental variable for `educ`.

We can do it in three ways: ILS, manual 2SLS, and full 2SLS

ILS

First stage:

```
fs = lm_robust(educ ~ nearc4 + demog + family, data=prox)
```

Reduced form:

```
rf = lm_robust(lwage ~ nearc4 + demog + family, data = prox)
```

ILS estimate:

```
rfcoef['nearc4']/fscoef['nearc4']
```

Manual 2SLS

Second stage: use `fitted.values`

```
ss = lm_robust(lwage ~ fs$fitted.values + demog + family, data=prox)
```

But standard errors are not valid.

Automated 2SLS

```
iv_robust(lwage ~ educ + demog + family | nearc4 + demog + family,  
data = prox)
```